




Scaling Users Shouldn't Mean Scaling Headcount

Introducing Voxe. AI infrastructure that breaks
the product-led support paradox.

Product-Led Growth Exposes the Limits of Linear Support


- Exponential **user growth** outpaces **linear headcount capabilities**.
- 52% of users who **encounter friction during onboarding** abandon within the first session.
- Silent churn is the **penalty** for **failing the speed test**.



The Fragmented Stack is Where Customer Context Goes to Die

A user chats on the website, emails later, and checks docs. Three interactions. Three separate tool silos. Zero shared memory.

Support agents lose 6 hours a week context-switching across 100+ SaaS apps.



Legacy AI Pricing Penalizes You for Improving

Incumbents charge
~**\$0.99** for every ticket
the AI resolves.

As you train the AI and
improve your
documentation, your
resolution rate climbs.

Your reward for
optimization is an
uncapped invoice.



AI Performance /
Resolution Rate




Monthly Invoice
Cost

40,000 AI resolutions = \$39,600/month.
Before any human seat costs.

A Unified Context Stream with Flat-Rate Economics

- Voxxe is not a chatbot wrapper. It is complete AI infrastructure.
- Coverage-based pricing means costs stay flat while performance scales.
- One system handling routing, retrieval, escalation, and human handoff.




From a Single URL to a Live Operation in 5 Minutes

- No manual FAQ writing or decision trees required.
- The system crawls, extracts, and structures your existing assets.




Smart Page Context Eliminates Generic Answers

- Most bots are context-blind. Voxe knows exactly where the user is in their journey.
- Reduces “How do I get started?” tickets by delivering the exact next step.




Beyond Static Docs with the Model Context Protocol (MCP)

- Answering from a PDF is tier-1 support. Real operations require live business logic.
- Voxe connects directly to your internal databases, proprietary APIs, and live systems.
- The AI answers based on what is true right now, not what was documented last month.




The Calendar Pipeline Converts Visitors into Booked Demos

- A sub-5-minute response time makes leads 9x more likely to convert.
- The AI checks real Google Calendar availability, presents slots, and confirms the meeting.
- No scheduling forms. No waiting for a human rep.



Omnichannel Routing Preserves the Complete Customer History

- 76% of customers expect consistent context across interactions.
- Multichannel traps data in silos. Omnichannel weaves it into a single timeline.
- A customer can start on WhatsApp and finish via email without repeating themselves once.



The Perfect Escalation Sequence



Semantic Routing: The AI routes to the correct team based on context, not just keywords.



Holding AI: Monitors liveness. If a human takes too long, the AI keeps the customer engaged.



Fusion & NeuroSwitch: Underlying control layers manage model routing, cost tracking, and quality filtering.




The Architectural and Economic Reality

	Legacy Stack (Intercom/Zendesk)	Voxe Infrastructure
Pricing Model	\$0.99 Per-Resolution (Penalizes success)	Flat-Rate Coverage (Scales efficiently)
Setup Time	Weeks of manual configuration	2-5 minutes via automated URL ingestion
System Access	Read-only from static PDFs	Read/Write to live databases via MCP
Human Handoff	Context frequently lost between tools	Full conversation continuity guaranteed

Track The Five Metrics That Actually Matter

Deflection rate is a vanity metric. Containment rate measures actual resolution.
AI First Response Time drops from the 15-minute human average to 23 seconds.



The 72-Hour Offline Experiment

A founder handed all support to Voxe and went fully offline for a 3-day weekend.
The system accurately retrieved answers, routed exceptions, and maintained brand tone.

34266

Total Conversations

Resolved Instantly by AI
(Avg response: 3.8s)

Semantically Routed to Human
(Waiting for Monday
with full context)

The 2 failures? Outdated documentation. Knowledge Gap Rate flagged them for immediate updates.

Stop Paying Per Resolution. Start Owning Your Operations.

Growth should not dictate your support costs. The infrastructure is ready.

[yourwebsite.com]

Deploy in 5 Minutes.

